

AWS FPGA Starter's Guide (Rev 0)

Introduction

This guide explains how to work with AWS FPGA Hardware/Software Development Kit (HDK/SDK) and deploy your FPGA application on AWS. Developing FPGA application in cloud flows like this.

- 1) Write your custom logic (CL) that goes into FPGA in SystemVerilog.
- 2) RTL Simulate your CL for debugging and validation.
- 3) Create Amazon FPGA image (AFI).
- 4) Write your host program in C/C++.
- 5) Remote into AWS FPGA instance and load AFI into FPGA.
- 6) Compile and run the C/C++ application.



Figure 1. Programming model

Creating AWS Account

First, sign up for an AWS account with your UW email account at <https://aws.amazon.com>. You will have to enter a valid credit card info, but you will not be charged for anything (Uncle Bezos got you covered this time). Once the sign up is complete, go to <https://aws.amazon.com/education/awseducate> to apply for free AWS credit. You want to choose the option to have the free credit applied to the account you just created (Do not pick the "AWS Starter Account" option). Be sure to do this step first, because this process may take up to 2-3 days at most. You can check whether your free credit has been applied to AWS at Billing Dashboard:

<https://console.aws.amazon.com/billing/home#/>.

The screenshot shows the AWS Billing Dashboard 'Credits' page. It includes a sidebar with navigation links, a main content area with a security check and a 'Redeem' button, and a table of redeemed credits.

Expiration Date	Credit Name	Credits Used	Credits Remaining	Applicable Products
2019-09-30	EDU_ENG_FY2017_Q3_MAN9_100USD	\$21.52	\$78.48	See complete list

Total Amount of Credits Remaining: \$78.48

Writing Custom Logic

Custom Logic (CL) is the core of your FPGA application that actually does something inside FPGA. HDK provides the shell template (SH) which encapsulates CL and provides I/O interface to the host and peripherals.

1) Setting up HDK

- 1) Log into one of EE Linux machines where Vivado is installed (remotely or in person). Those machines are named 'linux-lab-001.ee.washington.edu' through 'linux-lab-040.ee.washington.edu'.
- 2) Run the following line. 'settings64.sh' script will allow you call 'vivado' from command line, and appending this line to '.bashrc' will make sure that this script is called whenever you open a bash terminal.

```
$ "source /homes/lab.apps/xilinx_vivado_v2017.1_sdx/SDx/2017.1.op/settings64.sh" >> ~/.bashrc
```

- 3) Type 'bash' to switch terminal.
- 4) Create a directory where you want to store your project and cd in there.
- 5) Clone git repository from aws-fpga. It should create a directory named "aws-fpga".

```
$ git clone https://github.com/aws/aws-fpga.git
```

- 6) cd into aws-fpga directory, and you will find hdk_setup.sh file.

```
ERRATA.md      README.md      hdk            sdk
FAQs.md        RELEASE_NOTES.md hdk_setup.sh  sdk_setup.sh
LICENSE.txt    SDAccel       sdaccel_setup.sh
[dcjung@linux-lab-019 aws-fpga] $
```

Figure 2. inside aws-fpga directory

- 7) Run `source hdk_setup.sh`.

2) Hello World (Virtual LED/DIP switch)

Using the virtual LED and DIP switch is the most basic way for the host to communicate with FPGA. The host can pass some configuration or command via virtual DIP, and the FPGA can return the result or raise an alert to the host via virtual LED. It's called 'virtual', because we are remotely programming an FPGA board somewhere inside AWS datacenter, not a local FPGA board.

The example custom logics are located under 'aws-fpga/hdk/cl/examples/'. Those examples are provided by AWS, and you are free to take a look. Instead, we are going to use the 'cl_led_dip' custom logic in aws-fpga-starterpack.




Figure 3. aws-fpga-starterpack directory structure

- 1) Clone aws-fpga-starterpack repository in your project directory.

```
git clone https://github.com/tommydcjung/aws-fpga-starterpack.git
```

- 2) Copy the 'cl/cl_led_dip' folder from aws-fpga-starterpack to 'aws-fpga/hdk/cl/examples/'.

```
README.md      cl_examples_list.md  cl_hello_world_ref_hlx  cl_led_dip      hello_world_hlx
cl_dram_dma     cl_hello_world          cl_hello_world_vhdl     cl_uram_example
cl_dram_dma_hlx cl_hello_world_hlx      cl_ipi_cdma_test_hlx    common
[dcjung@linux-lab-019 examples]$
```

Figure 4. Inside aws-fpga/hdk/cl/examples/ directory

- 3) Inside the 'cl_led_dip' directory, you will find four directories.

```
[dcjung@linux-lab-019 cl_led_dip]$ ls
build  design  software  verif
```

'build' will be used for creating AFI later, and 'verif' will be used for simulating CL. All the custom logic Verilog code should be written in design.

- 4) cd into 'design', and you will find 'cl_led_dip.sv' file. This is the top-level Verilog module for your CL. As mentioned before, CL is encapsulated within shell (SH) to abstract I/O for the developer. The input and output ports of this top-level module are defined in the Verilog header file called 'cl_ports.vh', which is located in 'aws-fpga/hdk/common/shell_stable/design/interfaces/'.

```
// Hello World virtual DIP/LED
logic[15:0] vdip_q;
always_ff @(posedge clk_main_a0 or negedge rst_main_n_sync)
    if (!rst_main_n_sync) begin
        vdip_q <= 16'b0;
    end
    else begin
        vdip_q <= sh_cl_status_vdip[15:0];
    end
assign cl_sh_status_vled = {vdip_q[3:0], vdip_q[7:4], vdip_q[11:8], vdip_q[15:12]};
```

Figure 5. core functionality of cl_led_dip

The ports for virtual DIP and LED are called "sh_cl_status_vdip" and "cl_sh_status_vled" respectively, and they are each 16-bit wide. In the code snippet above, the input "sh_cl_status_vdip" flopped into the register called "vdip_q" every clock cycle, and the "cl_sh_status_vled" is assigned the reversed hex of "vdip_q".

Validating and Debugging CL by RTL simulation

It is a good idea to simulate your custom logic before you build and create the AFI. This provides a chance to validate that your design works as you intended. If there is a bug in your logic, this can be an invaluable tool for debugging. We will write the simulation code in SystemVerilog.

- 1) Go to 'verif/tests/' directory. There is a file called 'test_led_dip.sv'.

```
[dcjung@linux-lab-019 cl_led_dip]$ cd verif/
[dcjung@linux-lab-019 verif]$ ls
scripts  sim  tests
[dcjung@linux-lab-019 verif]$ cd tests/
[dcjung@linux-lab-019 tests]$ ls
test_led_dip.sv
```

- 2) Open 'test_led_dip.sv', and you will see that the test is setting the DIP switch to "0x0000", and then to "0xbeef". When you read the LED value, it should be "0xfeeb". More details about the test bench can be found here: https://github.com/aws/aws-fpga/blob/master/hdk/docs/RTL_Simulating_CL_Designs.md

```
module test_led_dip();

    import tb_type_defines_pkg::*;

    logic [15:0] vdip_value;
    logic [15:0] vled_value;

    initial begin

        tb.power_up();

        // set vDIP to 0x0000.
        tb.set_virtual_dip_switch(.dip(16'h0000));
        vdip_value = tb.get_virtual_dip_switch();
        $display("[%t] vdip = 0x%x", $realtime, vdip_value);

        vled_value = tb.get_virtual_led();
        $display("[%t] vled = 0x%x", $realtime, vled_value);

        // set vDIP to 0xbeef.
        tb.set_virtual_dip_switch(.dip(16'hbeef));
        vdip_value = tb.get_virtual_dip_switch();
        $display("[%t] vdip = 0x%x", $realtime, vdip_value);
        #10ns; // wait 10ns

        vled_value = tb.get_virtual_led();
        $display("[%t] vled = 0x%x", $realtime, vled_value);

        if (vled_value == 16'hfeeb)
            $display("[%t] Test passed.", $realtime);
        else
            $display("[%t] Test failed.", $realtime);

        tb.kernel_reset();
        tb.power_down();

        $finish;
    end
endmodule
~
~
```

Figure 6. test_led_dip.sv

- 3) In order to run the test, cd into "verif/scripts". There are three files of our interests.
 - a. Makefile – we are using this file to build and run the test.
 - b. top.vivado.f – if you add create a new module or include a new IP, make sure to include it in this file.
 - c. waves.tcl – during the simulation you can capture waveforms of some of the signals in your custom logic. The first line of 'wave.tcl' is "add_wave /card/fpga/CL". This statement will tell the simulation to capture all the signal in the top-level module (in this case 'cl_led_dip'). Note that it will

only signals at the top-level, but not the signals at a lower level. For example, if there was an instance named “MY_MODULE”, you will have to add ‘add_wave /card/fpga/CL/MY_MODULE’ to capture the signals inside “MY_MODULE” instance. You can also use “add_wave -recursive /card/fpga/CL” to capture all the signals at lower levels recursively, but it could result in a very large waveform output, so use with discretion.

- 4) While you are in “verif/scripts”, run “`make TEST=test_led_dip`”. The parameter after “TEST=” should match the name of the test file under “verif/tests” without the file extension (.sv). The test script will compile and run the test program, and it should print the output like this in the end.

```
[          5050000] vdip = 0x0000
[          5050000] vled = 0x0000
[          5050000] vdip = 0xbeef
[          5060000] vled = 0xfeeb
[          5060000] Test passed.
$finish called at time : 5160 ns : File "/home/dcjung/tutorial/aws-fpga/hdk/cl/examples/c
l_led_dip/verif/tests/test_led_dip.sv" Line 37
## quit
INFO: [Common 17-206] Exiting xsim at Thu Nov 16 06:57:21 2017...
bash-4.2$
```

Figure 7. test_led_dip test result

- 5) Now, let us view the waveform captured during the simulation. If you go back to “verif” directory, you will notice that “verif/sim” directory is created. The simulation results are stored under this directory, with the name of the test you just ran. cd into “verif/sim/test_led_dip”. There is a file called “tb.wdb”. That is the waveform database file. Create a text file called “open_waves.tcl” inside this directory, and add the following two lines and save.

```
current_filesset
open_wave_database tb.wdb
```

Now, run `vivado -source open_waves.tcl`

- 6) Under ‘Scope’ tab, you will see the hierarchy of modules. Click down to “/tb/card/fpga/CL”. In ‘Objects’ panel, you will see the signals defined under each hierarchy. Right-click ‘clk_main_a0’, and click “Add to Wave Window”. Also, add ‘rst_main_n’, ‘sh_cl_status_vdip[15:0]’ and ‘cl_sh_status_vled[15:0]’. Drag and click the waveform window to zoom into where the interesting events are happening.

You should see that virtual DIP starts out as ‘0x0000’ and becomes ‘0xbeef’. Then, one clock cycle after, virtual LED value becomes ‘0xfeeb’.

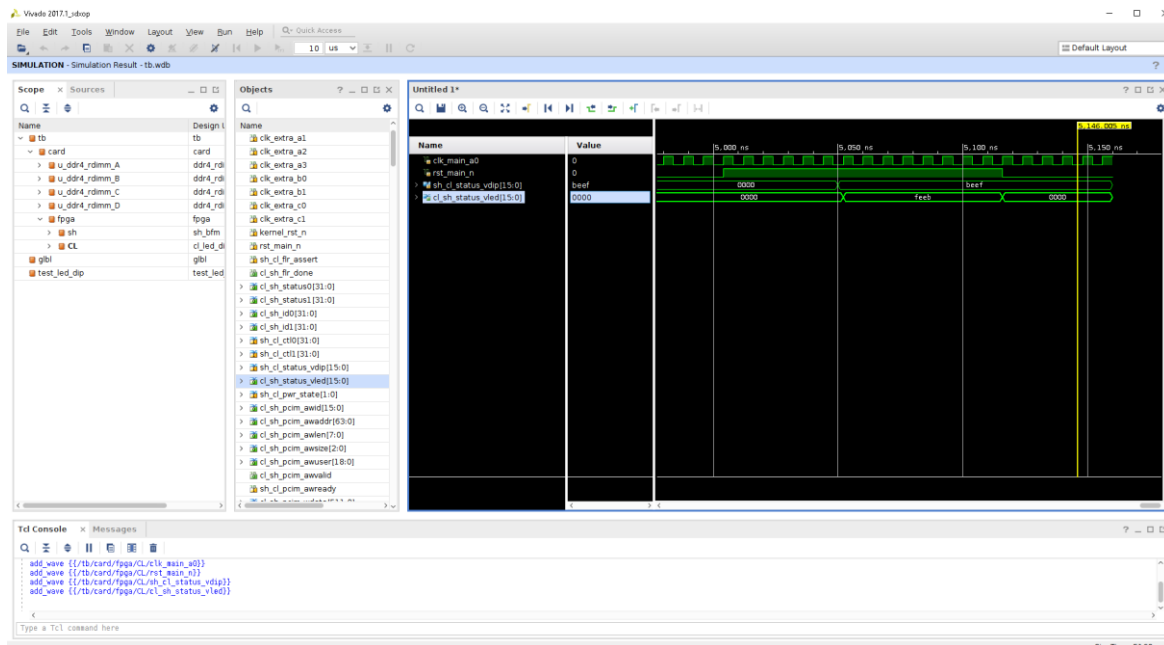


Figure 8. Waveform of `cl_led_dip`.

Creating Amazon FPGA Image

Now that we have validated the functionality of `cl_led_dip` through RTL simulation, it is time to create Amazon FPGA Image (AFI). First, we have to build a tarball from our CL design. Then, we have to run various AWS CLI commands to create AFI.

1. Creating tarball from CL design

- 1) cd into "aws-fpga/hdk/cl/examples/cl_led_dip"
- 2) Set the environment variable. When you type `echo $CL_DIR`, it should print out the path to your `cl_led_dip` directory.

```
$ export CL_DIR=$(pwd)
```

- 3) cd into "cl_led_dip/build/scripts" and run `./aws_build_dcp_from_cl.tcl -foreground`
- 4) The build process will take a long time even for the simplest application (~2-3 hours), so you can either move onto "Writing C/C++ Host Application" or go out for a walk and grab some coffee.
- 5) If the build was successful, cd into "cl_led_dip/build/checkpoints/to_aws/". You shall find a tarball file generated with timestamp as prefix (e.g. `17_11_16-183421.Developer_CL.tar`).

```
[dcjung@linuxsrv01 to_aws]$ pwd
/home/dcjung/tutorial/aws-fpga/hdk/cl/examples/cl_led_dip/build/checkpoints/to_aws
[dcjung@linuxsrv01 to_aws]$ ls
17_11_16-183421.Developer_CL.tar 17_11_16-183421.SH_CL_routed.dcp 17_11_16-183421.manifest.txt
[dcjung@linuxsrv01 to_aws]$
```

2. Configuring AWS CLI / Creating AFI

- 1) We are going to use AWS Command Line Interface (CLI) to create AFI. If you are working from your own computer you will have to install AWS CLI. <https://aws.amazon.com/cli/>
If you are working with EE dept machines, AWS CLI should have been installed.

- 2) We have to configure the AWS CLI so that it works with your AWS account credential. In order to do that we need to create access key pair. https://console.aws.amazon.com/iam/home?#/security_credential Click Access keys-> Create New Access Key. It will give you "Access Key ID" and "Secret Access Key". Save those two strings somewhere secure. Do not share the secret access key with others.

Search IAM

Dashboard
Groups
Users
Roles
Policies
Identity providers
Account settings
Credential report
Encryption keys

Your Security Credentials

Use this page to manage the credentials for your AWS account. To manage credentials for AWS Identity and Access Management (IAM) users, use the [IAM Console](#).

To learn more about the types of AWS credentials and how they're used, see [AWS Security Credentials](#) in AWS General Reference.

- + Password
- + Multi-factor authentication (MFA)
- Access keys (access key ID and secret access key)

You use access keys to sign programmatic requests to AWS services. To learn how to sign requests using your access keys, see the [signing documentation](#). For your protection, store your access keys securely and do not share them. In addition, AWS recommends that you rotate your access keys every 90 days.

Note: You can have a maximum of two access keys (active or inactive) at a time.

Created	Deleted	Access Key ID	Last Used	Last Used Region	Last Used Service	Status	Actions
Oct 15th 2017		[REDACTED]	2017-10-15T10:10:10Z	us-west-2	ec2	Active	Make Inactive Delete
Oct 21st 2017	Nov 16th 2017	[REDACTED]	N/A	N/A	N/A	Deleted	
Oct 6th 2017	Oct 21st 2017	[REDACTED]	N/A	N/A	N/A	Deleted	

[Create New Access Key](#)

- 3) Now back to the terminal. Run the following command. Enter the access key ID and secret access key. Default region name should be 'us-west-2' (Oregon). Just hit enter for default output format.

```
$ aws configure
```

```
[dcjung@linuxserv01 ~]$ aws configure
AWS Access Key ID [None]: [REDACTED]
AWS Secret Access Key [None]: [REDACTED]
Default region name [None]: us-west-2
Default output format [None]:
[dcjung@linuxserv01 ~]$
```

- 4) AWS S3 is a cloud-based file storage service. We have to upload our tarball to S3, which then will be picked up by AFI generation service to be processed.

- 5) Create a S3 bucket (analogous to a file directory)

```
$ aws s3 mb s3://my-fpga-cl --region us-west-2
```

- 6) Upload the tarball to the bucket just created.

```
$ aws s3 cp ./17_11_16-183421.Developer_CL.tar s3://my-fpga-cl/
```

- 7) Create a S3 bucket for AFI generation log.

```
$ aws s3 mb s3://my-fpga-log --region us-west-2
```

You should be able to verify that bucket has been created and the tarball uploaded as expected here on the AWS management portal: <https://s3.console.aws.amazon.com/s3/home?region=us-west-2#>

- 8) Start AFI generation.

```
$ aws ec2 create-fpga-image --region us-west-2 --input-storage-location Bucket=my-fpga-cl,Key=17_11_16-183421.Developer_CL.tar --logs-storage-location Bucket=my-fpga-log,Key=
```

Calling this command will return json output like this.

```
[dcjung@linuxsrv01 to_aws]$ aws ec2 create-fpga-image --region us-west-2 --input-storage-location Bucket=my-fpga-cl,Key=17_11_16-183421.Developer_CL.tar --logs-storage-location Bucket=my-fpga-log,Key=/
{
  "FpgaImageId": "afi-02c0bce69ef7c9902",
  "FpgaImageGlobalId": "agfi-0edfc64f852051323"
}
```

You will need to save these IDs somewhere to load AFI to FPGA later.

- 9) Check if AFI generation has completed. Once it's completed `FpgaImages.State.Code` will eventually become "available". Only then you are able to load this AFI to FPGA. It takes about 20~30 minutes.

```
[dcjung@linuxsrv01 to_aws]$ aws ec2 describe-fpga-images --fpga-image-ids afi-02c0bce69ef7c9902
{
  "FpgaImages": [
    {
      "UpdateTime": "2017-11-16T21:50:52.000Z",
      "FpgaImageGlobalId": "agfi-0edfc64f852051323",
      "State": {
        "Code": "pending"
      },
      "FpgaImageId": "afi-02c0bce69ef7c9902",
      "OwnerId": "240095434820",
      "CreateTime": "2017-11-16T21:50:52.000Z"
    }
  ]
}
```

Writing C/C++ Host Application

Developing host application has to be done on AWS virtual machine that runs FPGA Developer AMI (Amazon Machine Images). AWS provides C library for interacting with FPGA. This interaction is done through PCI Express Bus. Shell (SH) that encapsulates our CL handles the communication. This communication is generally divided into two physical function (PF): MgmtPF and AppPF. MgmtPF deals with the management functionality of FPGA: loading/unloading FPGA image, peeking/poking virtual DIP switch and LED. AppPF deals the functionalities that are more specific to your CL, such as sending and receiving stream of data over DMA.

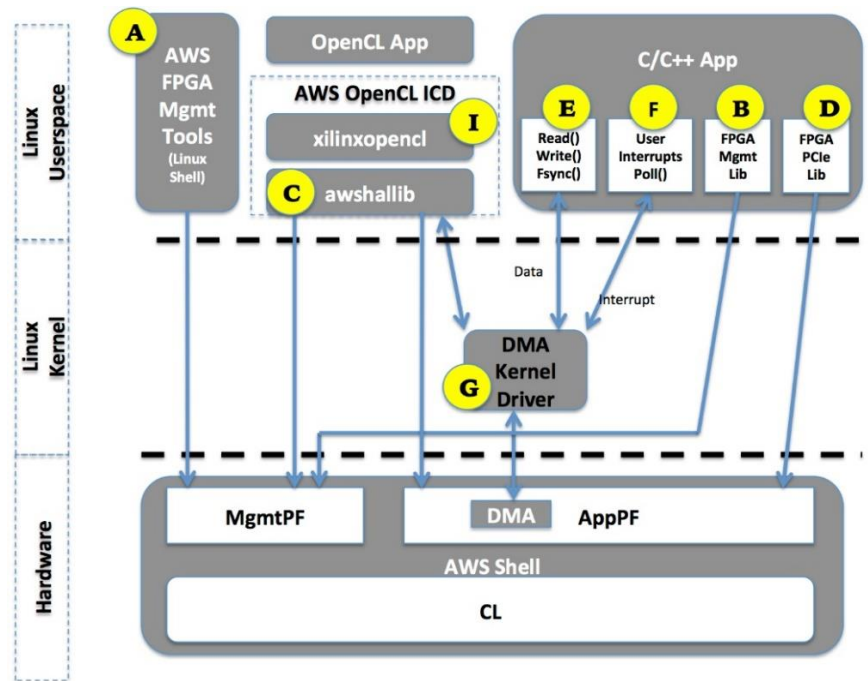
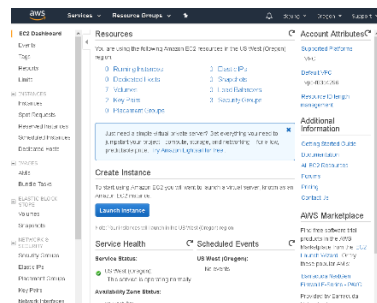


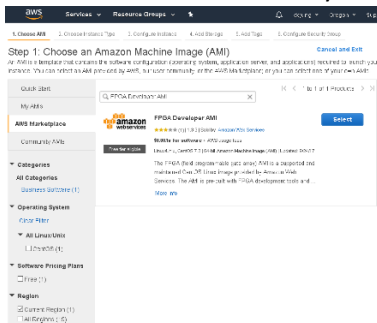
Figure 9. There are many ways to interact with FPGA.

1. Setting up FPGA Development AMI VM

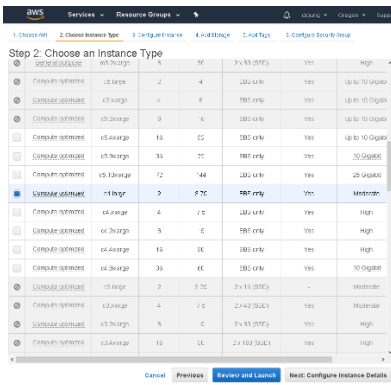
- 1) Go to <https://us-west-2.console.aws.amazon.com/ec2/v2/home?region=us-west-2>
Click “Launch Instance”.



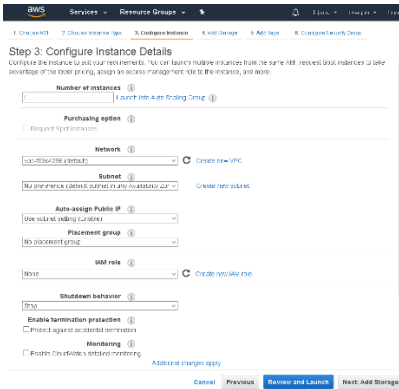
- 2) Click “AWS Marketplace” on the left tab, and search for “FPGA Developer AMI”. Click “Select”. Click “Continue” when it shows you a menu.




- 3) Select “c4.large”, and click “Next”.



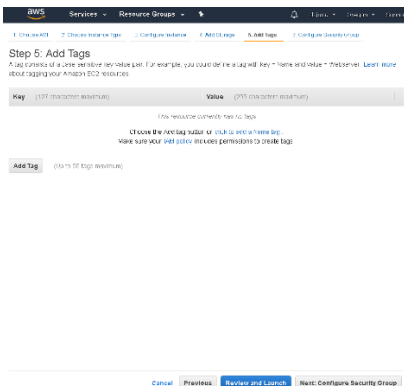
4) Leave the settings as they are, and click “Next”.



5) Leave the storage settings as they are, and click “Next”



6) Click “Next”



- 7) Click “Add Rule”. Select type “RDP” for remote desktop. Select source “Anywhere”. Click “Review and Launch”

Step 6: Configure Security Group

A security group is a virtual firewall that controls the traffic for your instance. On this page, you can add rules to allow specific traffic to reach your instance. For example, if you want to set up a web server and allow Internet traffic to reach your instance, you can add a rule that allows Internet access to the port 80 and 443 (HTTP) ports. You can create a new security group or select from an existing one below. Learn more about Amazon EC2 security groups.

Assign a security group: ☒ Create a new security group ☐ Select an existing security group

Security group name:

Description:

Type	Protocol	Port Range	Source	Description
SSH	TCP	22	Anywhere	SSH for admins
RDP	TCP	3389	Anywhere	RDP for admins

[Add Rule](#)

Warning
Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. The recommended best practice is to restrict access from specific IP addresses only.

[Cancel](#) [Previous](#) [Review and Launch](#)

- 8) Click “Launch”

Step 7: Review Instance Launch

AMI Details

PPGGA Developer AMI
PPGGA Developer AMI. Based on CentOS 7.3. Comes with Redis 5.0.10 installed.
Next Instance Type: [View Instance Type](#)

Hourly Software Fees: \$0.00 per hour on c4.large instance (Additional taxes may apply).
Software charges will begin once you launch this AMI and continue until you terminate the instance.
By launching this product, you will be subscribed to this software and agree that your use of this software is subject to the pricing terms and the seller's [End User License Agreement](#).

Instance Type

Instance Type	EC2	vCPUs	Memory (GB)	Instance Storage (GB)	EBX-Optimized Available	Network Performance
c4.large	8	2	3.75	EBX only	Yes	Moderate

Security Groups

Security group name:

Description:

Type	Protocol	Port Range	Source	Description
SSH	TCP	22	0.0.0.0/0	
RDP	TCP	3389	0.0.0.0/0	
RDP	TCP	3389	:/0	

[Cancel](#) [Previous](#) [Launch](#)

- 9) Select “Create a new key pair”. Download your key pair (.pem). Click “Launch Instances”.

Select an existing key pair or create a new key pair

A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about [removing existing key pairs from a public AMI](#).

Create a new key pair

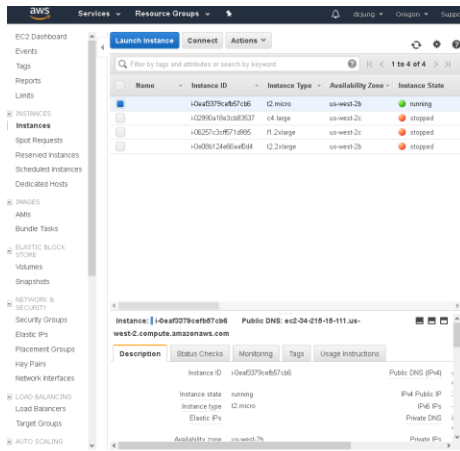
Key pair name:

[Download Key Pair](#)

You have to download the **private key file** (.pem file) before you can continue. Store it in a **secure and accessible location**. You will not be able to download the file again after it's created.

[Cancel](#) [Launch Instances](#)

- 10) You can see that your instance is up and running. Make sure that the instance state is “Stopped”, when you are not using. You will get billed for just leaving the machine running idly. You can click “Actions” dropdown bar, and select instance state.



2. Connecting to VM

- 1) You can follow the instruction here to connect to VM on Windows using PuTTY here:
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/putty.html?icmpid=docs_ec2_console
- 2) You can also connect using ssh:
<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AccessingInstancesLinux.html>

Note: your login will be 'centos'

3. Setting up SDK

- 1) Once you are logged into the VM, we need to set up the SDK. You want to store all your files in /home/centos/src/project_data, because the files stored there will be persist even after the session ends.

```
cd ~/src/project_data
```

- 2) Clone the aws-fpga repository.

```
git clone https://github.com/aws/aws-fpga.git
```

- 3) cd into aws-fpga, and run

```
source sdk_setup.sh
```

- 4) You should see message like this.

```
AWS FPGA: Copying Amazon FPGA Image (AFI) Management Tools to /usr/bin
AWS FPGA: Installing shared library to /usr/local/lib64
libfpga_mgmt.so.1 (libc6,x86_64) => /usr/local/lib64/libfpga_mgmt.so.1
AWS FPGA: Done with Amazon FPGA Image (AFI) Management Tools install.
Done with SDK install.
Done with AWS SDK setup.
[centos@ip-172-31-0-242 aws-fpga]$
```

4. Cloning aws-fpga-starterpack repository

I have created C++ wrapper around C library in the SDK.

- 1) cd back into ~/src/project_data
- 2) `git clone https://github.com/tommydcjung/aws-fpga-starterpack.git`
- 3) Inside the repository, there is aws-fpga-starterpack/host/common directory where C++ wrapper headers/classes are located.
- 4) You can go into aws-fpga-starterpack/host/cl_led_dip directory, where the source code for the host application for cl_led_dip is located. It also contains Makefile that builds the app.

```
#include <iostream>
#include "fabricmanager.h"
#include "fpga_mgmt.h"
using namespace std;

int main(int argc, char ** argv)
{
    auto fabricManager = new FabricManager();
    fpga_mgmt_image_info_t* info = 0;

    try
    {
        // get fpga image info.
        info = fabricManager->getImageInfo(0);
        cout << "FPGA Image Info:" << endl;
        printf("Vendor ID: 0x%x\r\n", info->spec.map[FPGA_APP_PF].vendor_id);
        printf("Device ID: 0x%x\r\n", info->spec.map[FPGA_APP_PF].device_id);

        // set virtual DIP to 0x0000
        fabricManager->setvDIP(0, 0x0000);
        cout << "Setting vDIP to 0x0000" << endl;

        // read virtual LED value
        uint16_t vLED = fabricManager->getvLED(0);
        printf("vLED = 0x%x\r\n", vLED);

        // set virtual DIP to 0xbeef
        fabricManager->setvDIP(0, 0xbeef);
        cout << "Setting vDIP to 0xbeef" << endl;

        // read virtual LED value
        vLED = fabricManager->getvLED(0);
        printf("vLED = 0x%x\r\n", vLED);
    }
    catch(exception& e)
    {
        cout << e.what() << endl;
    }

    // make sure to free dynamically allocated memory.
    free(info);
    delete fabricManager;

    return 0;
}
```

Figure 10. app.cpp

- 5) Run `make`. That should compile the program and create a file called `app`.

```
[centos@ip-172-31-0-242 cl_led_dip]$ ls
app.cpp Makefile
[centos@ip-172-31-0-242 cl_led_dip]$ make
g++ -o app ./app.cpp ../common/*.h ../common/*.hpp ../common/*.cpp -lfpga_mgmt -I/home/centos/src/project_data/aws-fpga/sdk/userspace/include/ -I../common/ -lrt -lpthread -std=c++0x
[centos@ip-172-31-0-242 cl_led_dip]$ ls
app app.cpp Makefile
[centos@ip-172-31-0-242 cl_led_dip]$
```

- 6) From here, you can develop your own FPGA host application in C++. You should create your own github repository and start from there.

Loading AFI in AWS FPGA instance

Finally, we have CL and host application written. We are ready to run our application on the real FPGA instance.

1. Setting up F1 instance

Amazon Elastic Compute Cloud (EC2) now provides F1 instances, which are equipped with Xilinx FPGA (VU9P). The details of this FPGA can be found under 'doc/' in aws-fpga-starterpack repository. (ds890-ultrascale-overview.pdf).

F1 Instance Details

Instance Type	FPGA Cards	vCPUs	Instance Memory (GiB)	SSD Storage (GB)	Enhanced Networking	EBS Optimized
f1.2xlarge	1	8	122	470	Yes	Yes
f1.16xlarge	8	64	976	4 x 940	Yes	Yes

For F1.16xlarge instances, the dedicated PCI-e fabric lets the FPGAs share the same memory space and communicate with each other across the fabric at up to 12 GBps in each direction.

1. The steps for setting up f1 instance is nearly identical to setting up FPGA Development AMI VM as above except we are picking f1.2xlarge as instance type. We will still be using FPGA Development AMI as the machine image.

Running C/C++ host application

1. Setting up aws-fpga and aws-fpga-starterpack repo
 - 1) Once you are logged into f1.2xlarge instance, do the same as done above. Go to ~/src/projectdata. Git clone both aws-fpga, and aws-fpga-starterpack repositories.
 - 2) Run `source sdk_setup.sh` from aws-fpga directory.
 - 3) Run `make` to build the host application from aws-fpga-starterpack/host/cl_led_dip.
 - 4) Clear the AFI slot 0.

```
$ sudo fpga-clear-local-image -S 0
```

```
[centos@ip-172-31-0-242 cl_led_dip]$ sudo fpga-clear-local-image -S 0
AFI          0          none          cleared          1          ok          0          0x071417d3
AFIDEVICE    0          0x1d0f          0x1042          0000:00:1d.0
[centos@ip-172-31-0-242 cl_led_dip]$
```

- 5) Load the AFI you created in the section “[Creating Amazon FPGA Image](#)” section above. Use the FPGA image global ID that started with ‘agfi-’ that you saved somewhere .

```
$ sudo fpga-load-local-image -S 0 -I {{ Your AFGI ID }}
```

```
[centos@ip-172-31-0-242 cl_led_dip]$ sudo fpga-load-local-image -S 0 -I agfi-0edfc64f852051323
AFI          0          agfi-0edfc64f852051323 loaded          0          ok          0          0x071417d3
AFIDEVICE    0          0x1d0f          0xf000          0000:00:1d.0
[centos@ip-172-31-0-242 cl_led_dip]$
```

- 6) Run the host application.

```
$ sudo ./app
```

```
[centos@ip-172-31-0-242 cl_led_dip]$ pwd
/home/centos/src/project_data/aws-fpga-starterpack/host/cl_led_dip
[centos@ip-172-31-0-242 cl_led_dip]$ ls
app.cpp  Makefile
[centos@ip-172-31-0-242 cl_led_dip]$ make
g++ -o app ./app.cpp ../common/*.h ../common/*.hpp ../common/*.cpp -lfpga_mgmt -I/home/centos/src/project_data/aws-fpga/sdk/userspace/include/ -I../common/ -lrt -lpthread -std=c++0x
[centos@ip-172-31-0-242 cl_led_dip]$ ls
app  app.cpp  Makefile
[centos@ip-172-31-0-242 cl_led_dip]$ sudo ./app
FPGA Image Info:
Vendor ID: 0x1d0f
Device ID: 0xf000
setting vDIP to 0x0000
vLED = 0x0
setting vDIP to 0xbeef
vLED = 0xfeeb
[centos@ip-172-31-0-242 cl_led_dip]$
```

Figure 11. It works.

- 7) Make sure that your instances are turned off when you are done.

Useful Links

<https://github.com/aws/aws-fpga>

<https://forums.aws.amazon.com/forum.jspa?forumID=243&start=0>

<https://github.com/tommydcjung/aws-fpga-starterpack>

<https://github.com/>

<https://aws.amazon.com/console/>

<https://aws.amazon.com/education/awsseduate/>